


How per-category research works

One `run_category` subtask, top to bottom. The four LLM steps each show their input and output; everything else is deterministic Python. The curator can loop back to search again, bounded.

- LLM call** — one `call_llm` (retry + non-fatal fallback)
- Deterministic code**
- Data store** (SQLite)



Fenced write = the atomic, generation-guarded persist that keeps refresh / reaper races correct. Source: `research/pipeline.py`, `research/tasks.py`.